


CBOW



Skip-gram

Word Representations

Jay Urbain, PhD

jay.urbain@gmail.com

<https://www.linkedin.com/in/jayurbain/>

EMBEDDINGS

WORKING WITH TEXT

- Neural networks don't take raw text as input. Only numeric tensors.
- *Vectorizing text* is the process of transforming text into numeric tensors. This can be done in multiple ways:
 - Segment text into words, and transform each word into a vector.
 - Segment text into characters, and transform each character into a vector.
 - Extract n-grams of words or characters, and transform each n-gram into a vector.
- Apply categorical transformation:
 - One-hot encoding
 - Word embedding

ONE-HOT ENCODING

- Words are categorical, they have no *ordinal* relationship. So they can't just be numerically encoded.
- For each word:
 - Numerically encode each word: {red:0, green:1, blue:2}
 - Generate vector of zeros the length of all possible words/categories
 - Set the index value for the word in the vector


red,	green,	blue
1,	0,	0
0,	1,	0
0,	0,	1

WORD EMBEDDINGS

- One-hot results in *sparse vector representation* (mostly zeros, hi dimensions).
- Word embeddings provide *dense vector representations*.
- Idea is to “embed” words into a lower-dimensionality space.
- The dimensions of this space are typically defined by word context, i.e., semantically similar words are embedded near each other.
- Popular algorithms:
 - PMI, Word2Vec: Skip-gram or CBOW, GLoVE, fasttext, ...



One-hot word vectors:
- Sparse
- High-dimensional
- Hardcoded



Word **embeddings**:
- Dense
- Lower-dimensional
- Learned from data

The Distributional Hypothesis (Firth, 1957)

- ‘You can tell a word by the company it keeps’
 - The *cat* sat on the mat
 - The *dog* sat on the mat
 - The *elephant* sat on the mat
 - The *quickly* sat on the mat

Embeddings: distributed representation of text

Text is represented *categorically*, but we want similar words to have similar representations

the	gray	cat	sat	on	the	gray	mat
0.4	1.0	0.9	0.1	0.5	0.4	0.9	0.2
0.1	0.1	1.0	0.9	0.3	0.1	1.0	1.0
0.5	0.1	0.1	0.2	0.8	0.5	0.1	0.1
0.7	0.5	0.2	0.3	0.9	0.7	0.2	0.1

```
embedding_layer = nn.Embedding(num_embeddings=vocab_size, embedding_dim=4)
embeddings = embedding_layer('the gray cat sat on the gray mat'.split())
```

Embeddings: Word2Vec

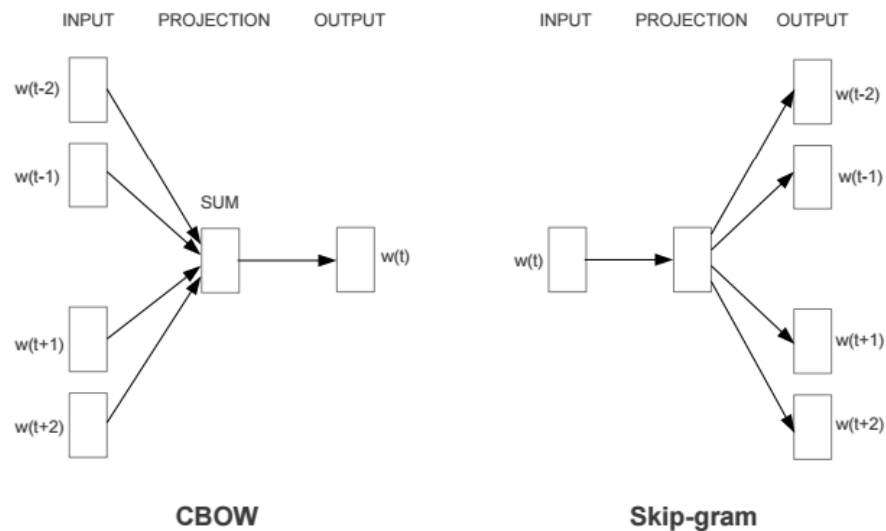
Idea: Maximize the probability of any context word given the current center word:

$$J'(\theta) = \prod_{t=1}^T \prod_{\substack{-m \leq j \leq m, \\ j \neq 0}} p(w_{t+j} | w_t; \theta)$$

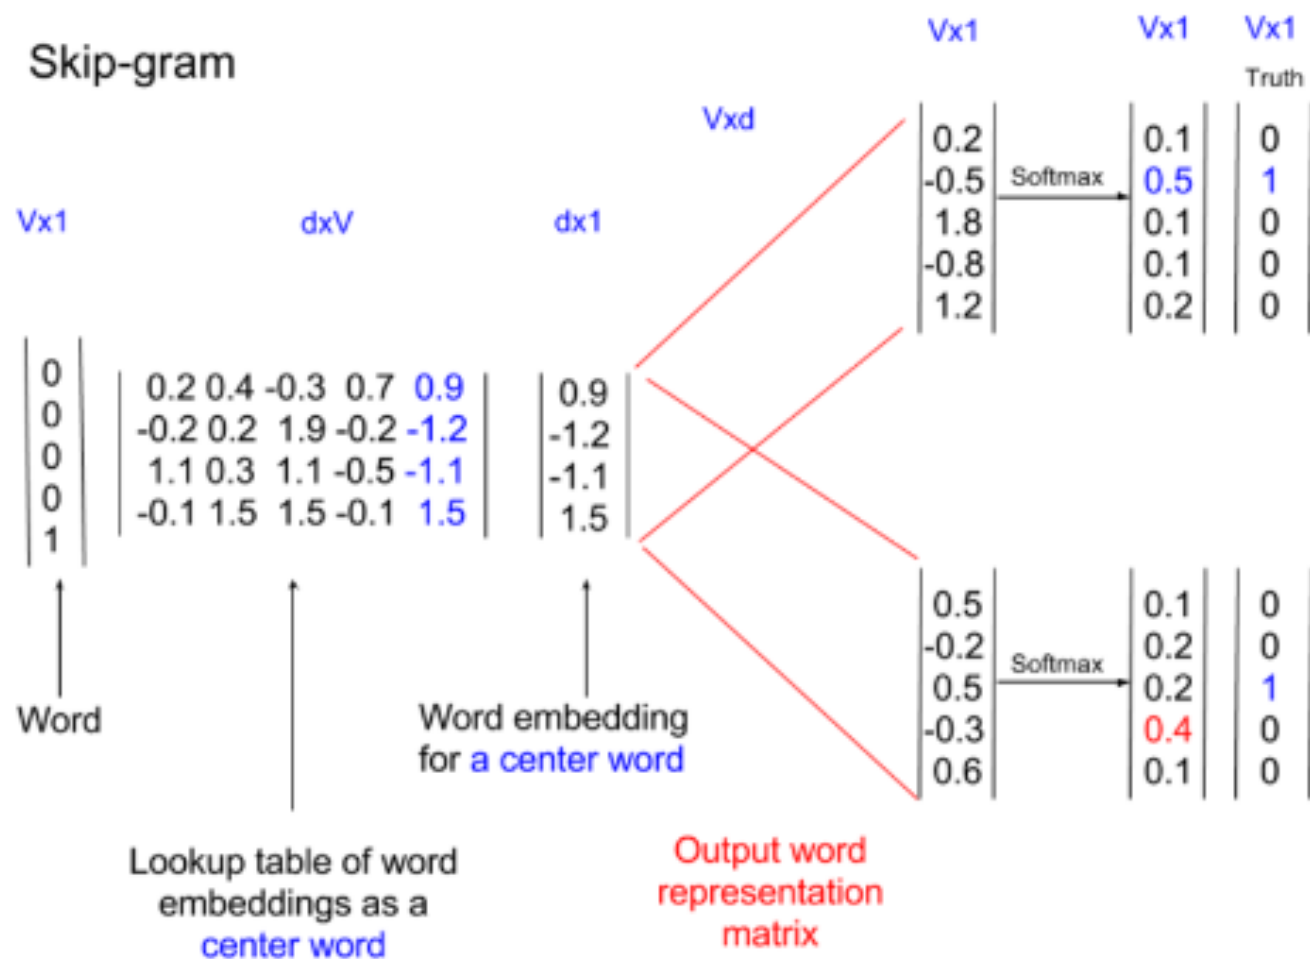
- ▶ T - our text(corpus). For each word $t=1 \dots T$, we try to predict surrounding words
- ▶ m - size of our window
- ▶ θ represents all variables we will optimize

Embeddings: Word2Vec

- **CBOW** model predict missing word (focus word) using context (surrounding words).
- **Skip gram** model predicts context based on the word in focus.
- Context is a fixed number of words to the left and right of the word in focus.
- Maximize average log probability of context words co-occurring with focus words.



Skip-gram



Pointwise Mutual Information

- Normalized skip gram probability
- Log probability of context word, x , co-occurring with target word y .

$$\text{pmi}(x; y) \equiv \log \frac{p(x, y)}{p(x)p(y)}$$

- Urbain, J., Bushnee, G., Knudson, Kowalski, G., P., Taylor, B. "Distributional Semantic Concept Models for Entity Relation Discovery," NAACL VSM-NLP (National Association of Computational Linguistics workshop on Vector Space Modeling in Natural Language Processing), June 2015.
- Levy, O., and Goldberg, Y., Neural Word Embedding as Implicit Matrix Factorization, NIPS, 2014.
- Urbain, J. A Distributed Dimensional Indexing Model for Information Retrieval and Extraction. NLM/NIBIB Workshop on Natural Language Processing: State of the Art, Future Directions and Applications for Enhancing Clinical Decision-Making, (April 23-24, 2012).

GloVe Word Embeddings

1. Collect word co-occurrence statistics in a word co-occurrence matrix X . Where each element X_{ij} represents how often word i appears in context of word j .
2. Define soft constraints foreach word pair. Word vectors $\sim \log$ of co-occurrence.

$$w_i^T w_j + b_i + b_j = \log(X_{ij})$$

3. Define a cost function to optimize.

$$J = \sum_{i=1}^V \sum_{j=1}^V f(X_{ij})(w_i^T w_j + b_i + b_j - \log X_{ij})^2$$

$f(X)$ is a weighting function to prevent learning only extremely common word pairs.

$$f(X_{ij}) = \begin{cases} (\frac{X_{ij}}{x_{max}})^\alpha & \text{if } X_{ij} < XMAX \\ 1 & \text{otherwise} \end{cases}$$

<https://nlp.stanford.edu/projects/glove/>

FastText

- Each word is represented as a bag of character n-grams.
- Words are represented as the sum of their character n-grams.
- Captures morphology of words.
- Provides a representation for rare and out-of-vocabulary words.
- Very fast.
- Outperforms Word2Vec and GloVe in most benchmarks.

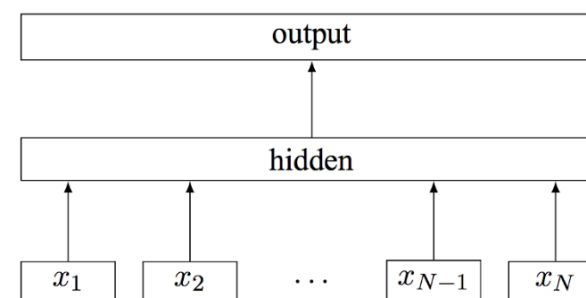


Figure 1: Model architecture of `fastText` for a sentence with N ngram features x_1, \dots, x_N . The features are embedded and averaged to form the hidden variable.

Image taken from the original paper

E.g., tri-grams for the word *apple* are $\{_ap, app, ppl, ple, le_ \}$.

Bojanowski, P., Grave, E., Joulin, A, Mikolov, T. Enhancing Word Vectors with Subword Information, 2017.

Selected word embedding: GloVe + Character Embedding

- Embedding for each word w is created by concatenating pretrained $p1=300$ GloVe vector (Pennington et al., 2014) and the word's character embedding.
- GloVe word vectors are fixed during training.
- All the out-of-vocabulary words are mapped to a token, whose embedding is trainable with random initialization.
- Adapted from [Seo et al., 2016](#), and [Srivastava et al., 2015](#)).

Problems with word embeddings?

Synonyms:

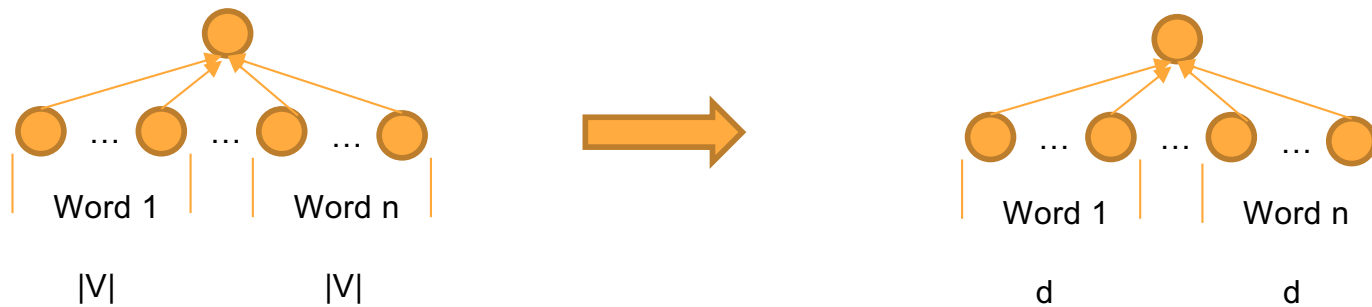
- River *bank*, *Bank* shot, *bank* deposit all have the same word representations.
- ... more on learning general models of language later.

Properties of embeddings

- They are dense: 0.53, 0.2, -1.2,
- They are low dimensional: (typically $50 \leq d \leq 300$)
- Embed domain semantics:
 - $\text{'king'} - \text{'man'} + \text{'woman'} \approx \text{'queen'}$
 - $\text{'paris'} - \text{'france'} + \text{'spain'} \approx \text{'madrid'}$
- Generalize easily!

Properties of embeddings (cont.)

- They are used as input across a variety NLP tasks (transfer learning!):




Total #
of input
parameters:


$$n * |V|$$

>>

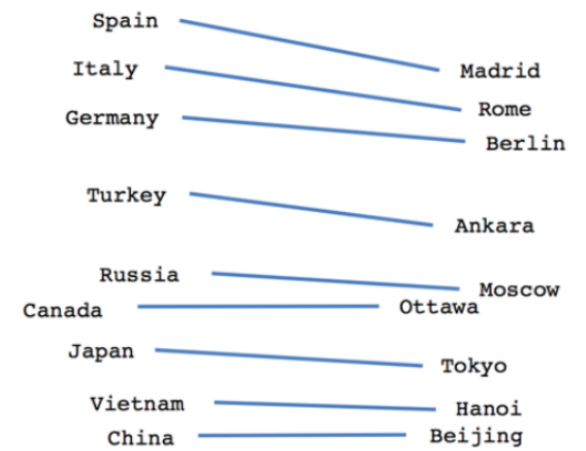
$$n * d$$



Male-Female



Verb tense



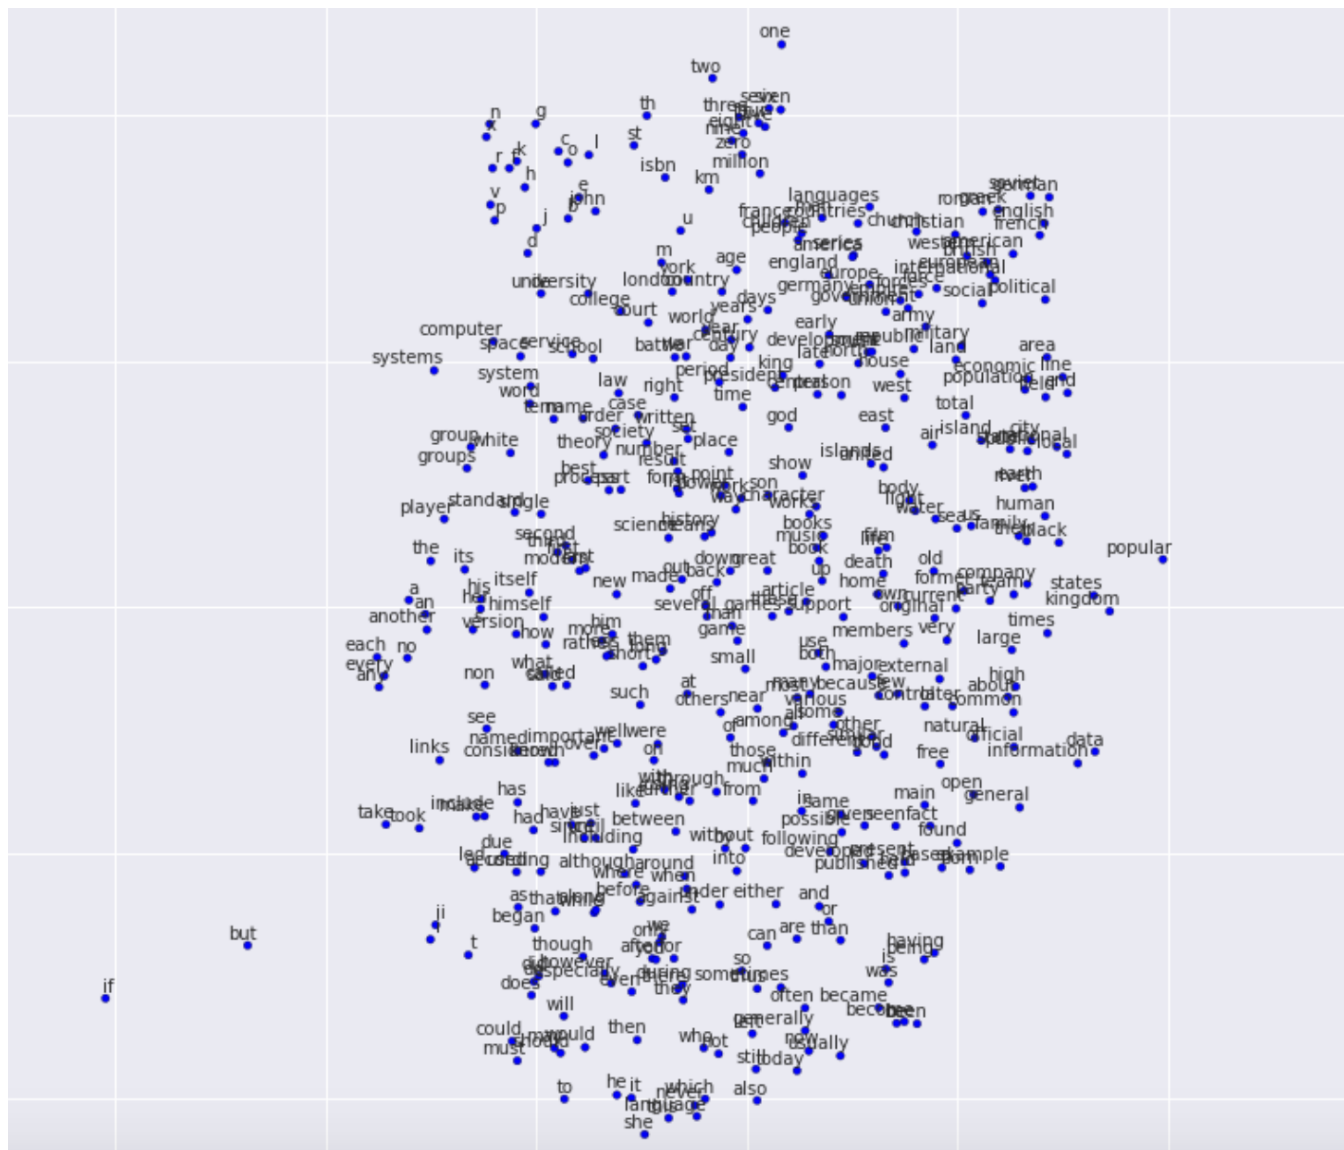
Country-Capital

<https://www.tensorflow.org/tutorials/representation/word2vec>

Word2Vec Galaxy

- https://www.youtube.com/watch?v=SnmQdKq_2hk

https://www.youtube.com/watch?v=SnmQdKq_2hk



USING WORD EMBEDDINGS

1. Learn an Embedding
 - Use embedding layer in your deep learning model
2. Reuse an Embedding
 - Download pre-trained word vectors

Lab 2 - Word Representations

https://github.com/jayurbain/DeepNLPIntro2019/blob/master/notebooks/embeddings/word_vector_visualization.ipynb